

## Concurrency control protocols:

Main aim to achieve serializability or recoverability  
 For this we use Concurrency control protocols  
 To Achieve this we use locking protocols.

### Shared - Exclusive Protocols:-

#### Shared Exclusive lock :-

Shared lock(s) :- if transaction locked data item in shared mode then allowed to read only.

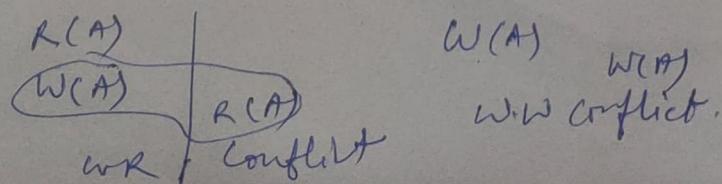
Exclusive lock (X) :- if transaction locked data item in exclusive mode then allowed to read & write both.

For $T_1$	$T_1$	$T_2$	
Shared lock $\rightarrow S(A)$		$X(A)$	we can't take shared lock because we have to read & write both
only read $\rightarrow R(A)$		$R(A)$	so we used exclusive lock.
Unlock $\rightarrow U(A)$		$W(A)$	

After Completion of work.

#### Compatibility of locks:-

		S Request $\rightarrow X$	
Grant S		Yes	No
X	Yes	Yes	No
	No	No	No



Grant - Shared lock on A request  $R(A) \rightarrow R(A)$   
 No conflict

$R(A)$  |  $R(A)$   
 $W(A)$  |  $W(A)$   
 Conflict may occur

## Problems in S/X locking :-

- 1) May not sufficient to produce only serializable schedule. — ~~Not enough~~
  - 2) May not free from Irrecoverability.
  - 3) May not free from deadlock. (Infinite waiting)
  - 4) May not free from starvation. (Finite waiting)
- Ex:-
- |                              |                              |
|------------------------------|------------------------------|
| T1                           | T2                           |
| X(A)<br>R(A)<br>W(A)<br>U(A) | X(B)<br>R(B)<br>W(B)<br>U(B) |
- when T1 unlock A.  
SCA → Not serializable
- cycle/loop
- 2PL (2 Phase locking)
- Ex:-
- |                              |                              |
|------------------------------|------------------------------|
| T1                           | T2                           |
| X(A)<br>R(A)<br>W(A)<br>U(A) | X(A)<br>R(A)<br>W(A)<br>U(A) |
- lock & commit  
unlock & release
- X(B) want  
→ X(A) w
- Two transactions  
wait for resource  
& nature is infinite loop.

- Growing Phase : locks are acquired and no locks are released
- Shrinking Phase - locks are released and no locks are acquired.

S/X
T1
X(A)
S(B)
R(A)
W(A)
R(B)
S(A)
R(C)
U(A)

Here we acquire locks.  
growing Phase

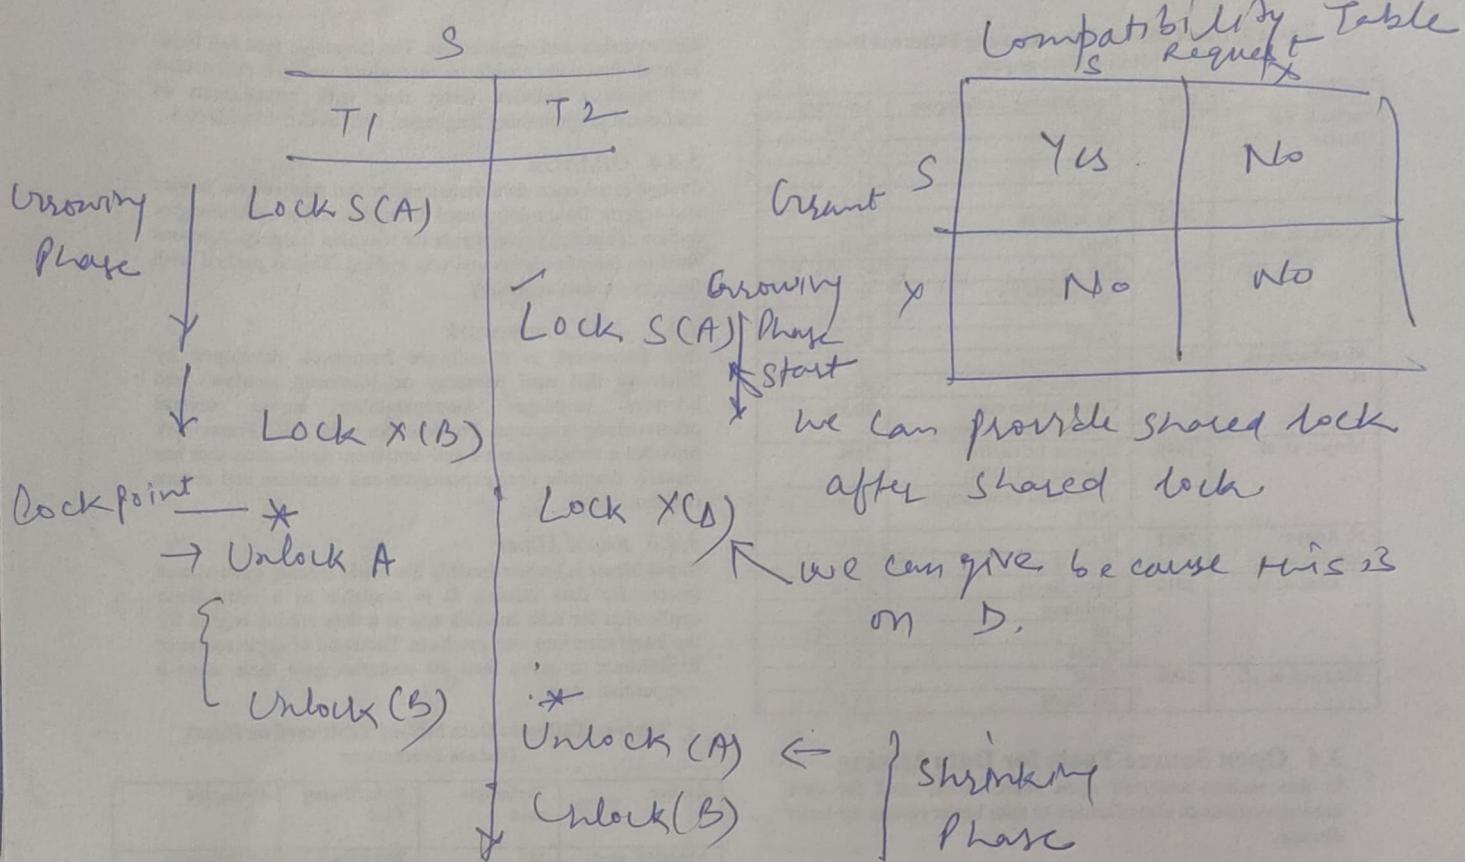
By this we achieve serializability.

T1	T2
X(A) R(A) W(A)	S(B) R(B) U(A)
	R(A) (wait)

don't enter  
till T2.

starts here After that we can't acquire locks

Transactions who are using 2PL that it is always serializable  $T_1 \rightarrow T_2$



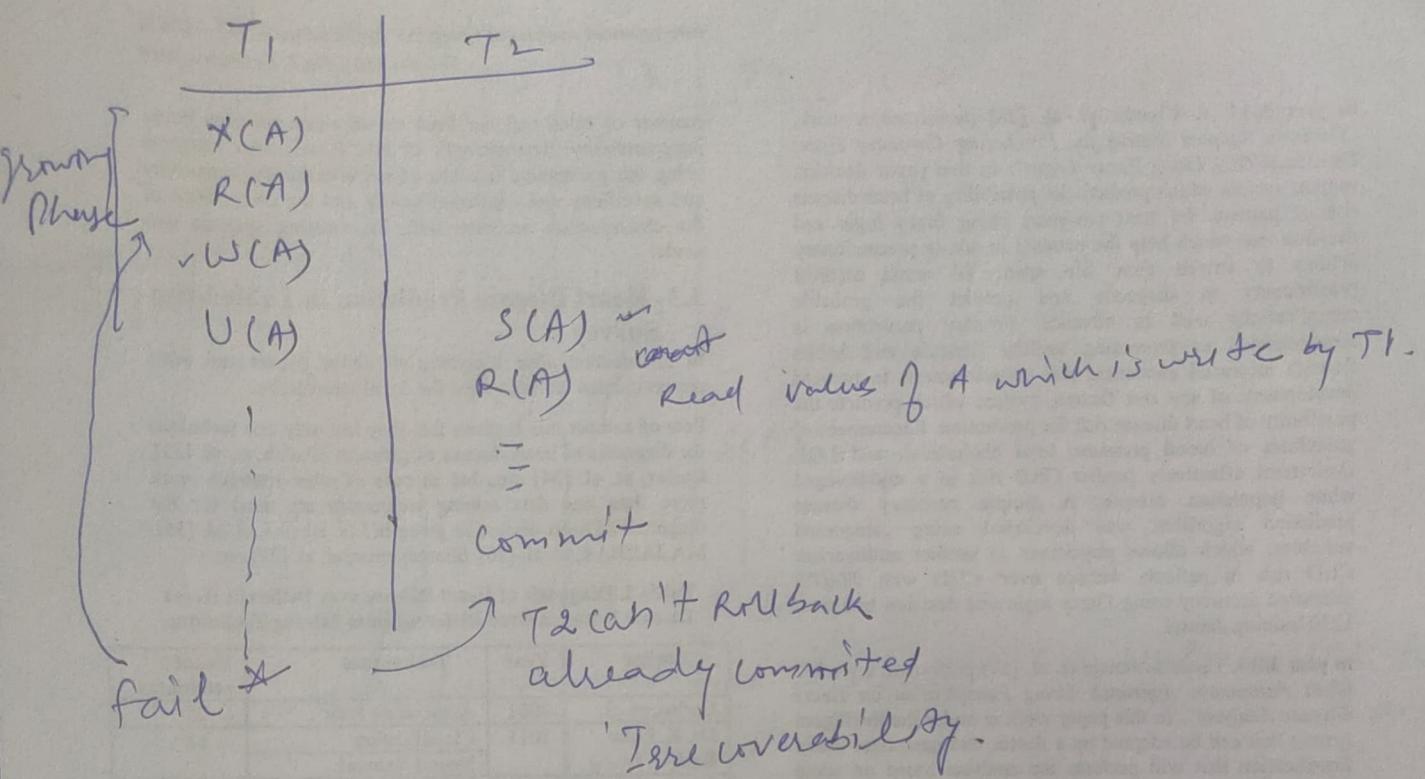
→ Lock Point - where the ~~leasing~~ transaction's first unlock starts that point is called lock point.

Problems in 2PL :- Advantage - Always ensures serializability.

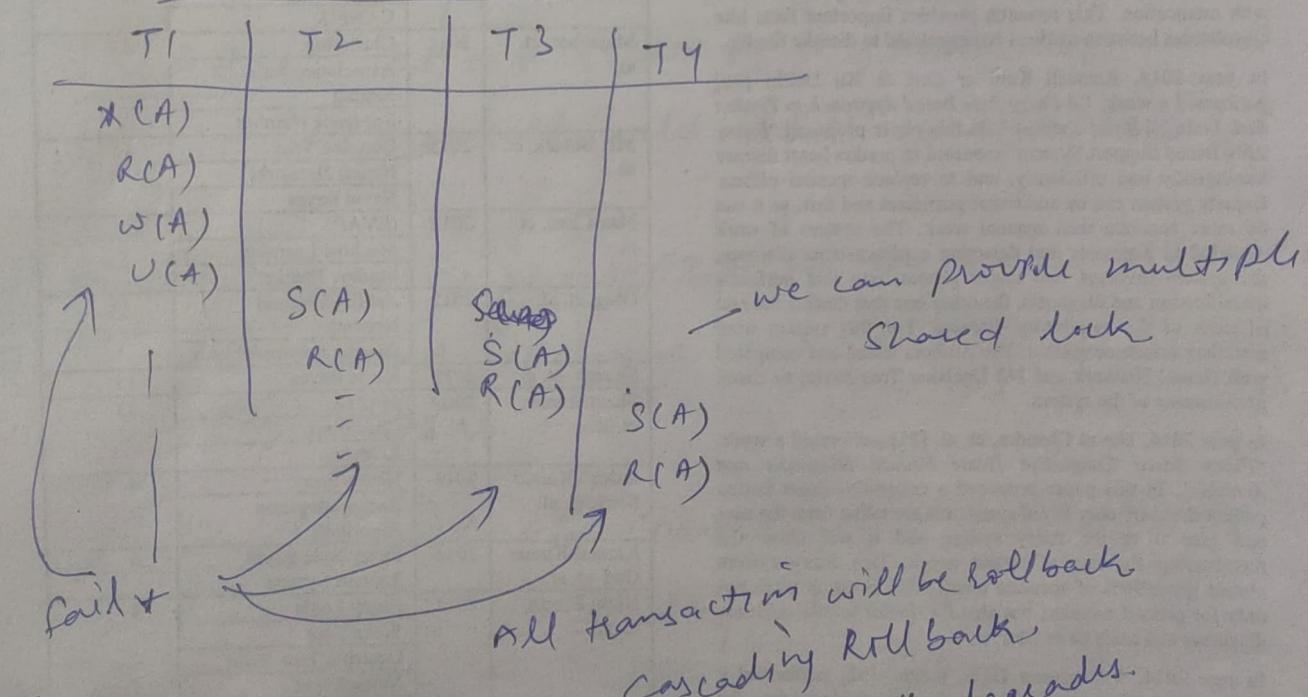
Disadvantages :- ① May not free from irrecoverability.

- ② Not free from deadlocks.
- ③ Not free from starvation.
- ④ Not free from cascading Rollback.

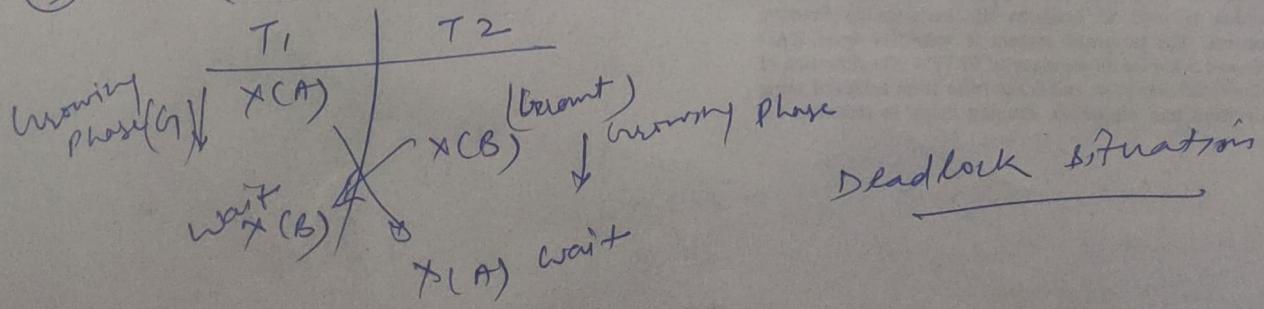
① may Not free from inconsistency.



② Free from cascading rollback:-



③ NOT free from Deadlocks :-



## (4) Not free from Starvation:

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>	T <sub>4</sub>
	S(A)		
Exclusive lock → X(A) Not possible			
		S(A)	
	U(A)		
		U(A)	
			S(A)
			U(A)

T<sub>1</sub> transaction has to wait till transaction T<sub>4</sub> unlock the Data item.  
It will take long time to wait.

### Extension of 2PL

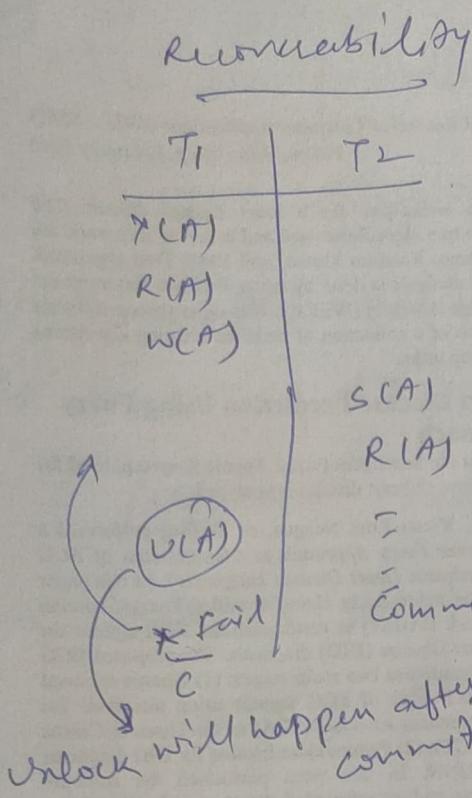
Strict 2PL: It should satisfy the basic 2PL and all exclusive locks should hold until Commit/Abort.

Rigorous 2PL: It should satisfy the basic 2PL → all shared, exclusive locks should hold until Commit/Abort.

This will resolve the problem of irrecoverability & Cascading Rollback.

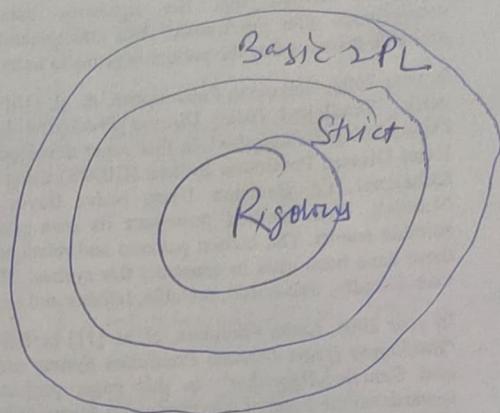
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
X(A)		
R(A)		
W(A)		
	S(A)	=
	R(A)	
		S(A) =
		R(A)

STRICT 2PL says: Exclusive locks should hold until Commit/Abort.  
Before Unlock we have to wait Commit/Abort.  
the transaction

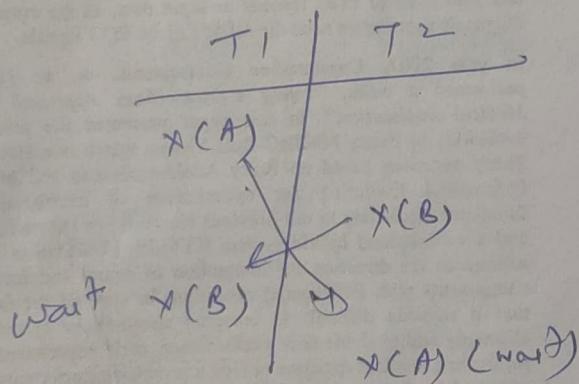


Advantages

- Cascadable
- Strict Recoverable



Disadvantage — Not free from deadlock + starvation

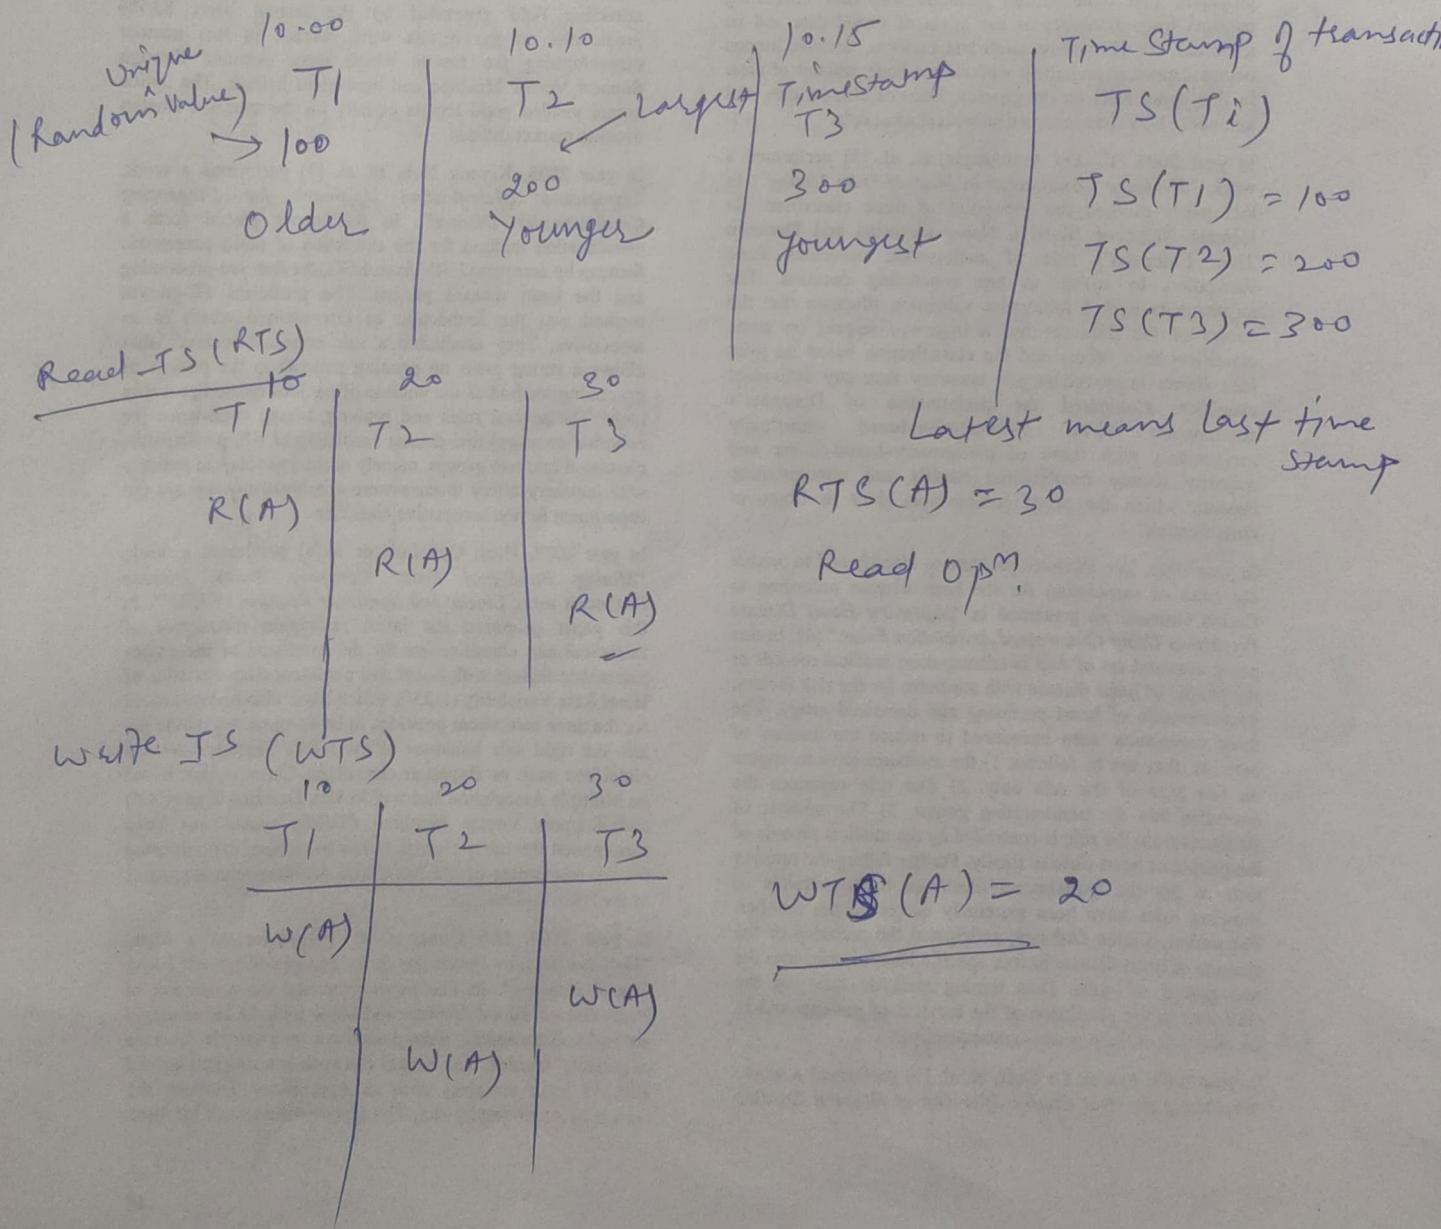


### ③ Conservative 2PL —

Before start any transaction it will occupy all the locks that Transaction T1 wants. For this prediction algorithm is to be used. It will solve the problem of deadlock. But in DBMS it is not possible to predefined what locks we require in start/advance.

## Time Stamp Ordering Protocol

- Time stamp is unique value assign to every transaction
- Tells the order (when they enter into system)
- $\text{Read-TS (RTS)} = \text{last (latest) transaction no. which performed read successfully.}$
- $\text{Write-TS (WTS)} = \text{last (latest) transaction no. which performed write successfully.}$



## Rules :-

① Transaction  $T_i$  issues a read(A) op<sup>n</sup>

a) if  $WTS(A) > TS(T_i)$ , rollback  $T_i$

b) otherwise execute R(A) operation

$$\text{Set } RTS(A) = \max \{ RTS(A), TS(T_i) \}$$

② Transaction  $T_i$  issues a write(A) op<sup>n</sup>

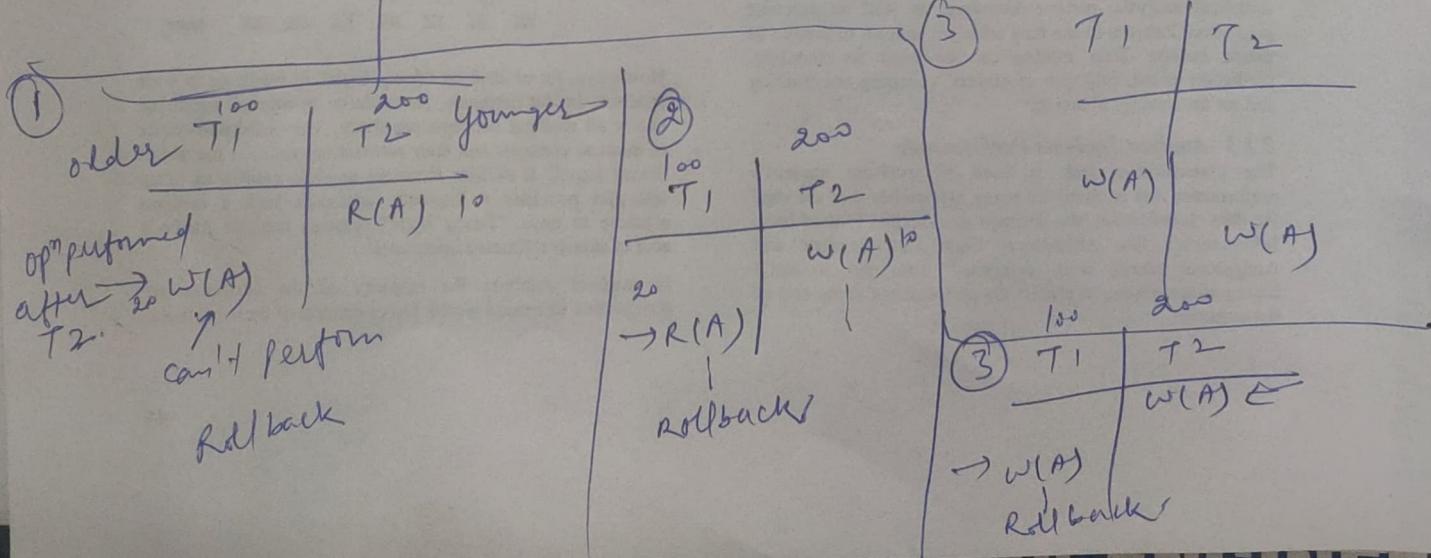
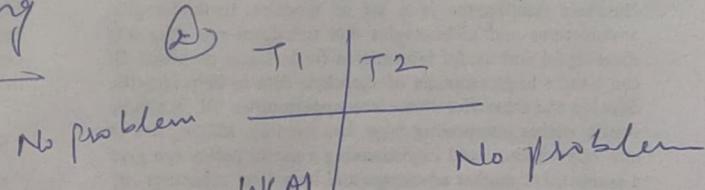
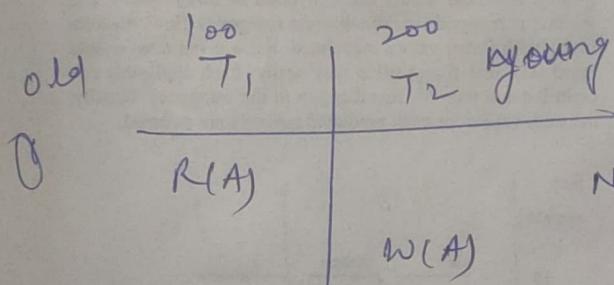
a) If  $RTS(A) > TS(T_i)$  then rollback  $T_i$

b) If  $WTS(A) > TS(T_i)$  then rollback  $T_i$

c) Otherwise execute write(A) operation

$$\text{Set } WTS(A) = TS(T_i)$$

Remember - Older transaction comes first, then first it will complete



Time Stamp Ordering Protocol

1st oldest will execute

$B \rightarrow L$ (100) oldest $T_1$	Time Stamp (200) $T_2$	Ordering (300) $T_3$ youngest	Protocol
$R(A)$			
$w(C)$	$R(B)$		
$R(C)$	$R(B)$		
$w(B)$	$w(A)$		
$wTS(A)$		$TS(T_1)$	
$0 > 100$ $R(A)$ — Rule 1(a)			Rule 1(a) $0 > 100 \rightarrow$ false move to otherwise condition
$0 > 200$ $R(B)$ — Rule 1(a)			(b) set $RTS(A) = \max(0, 100)$
$0 > 100 ] w(C)$ — Rule 2(a)			$0 > 200$ false $\rightarrow RTS(B) = \max(0, 200)$
$0 > 100 ] w(B)$ — Rule 2(a)			$TS(T_2) > TS(T_1)$ update $- 200$
$0 > 300 R(B)$			(c) $0 > 100$ set $wTS(C) = TS(T_2)$
$100 > 100 R(C)$ — false			update $C(0) \leftarrow 100$
$200 > 200 ] T_2 \xrightarrow{w(B)} R(B)$ balk			Restart after $T_3$ .
$100 > 300 ] w(A)$			Neglect all the values of $T_2$ transactions
$0 > 300 ]$			Restart after $T_3$ .

Initial zero-zero.

(2)

## Multiple Granularity Protocol :-

Granularity means hierarchy of the level or different levels of data.

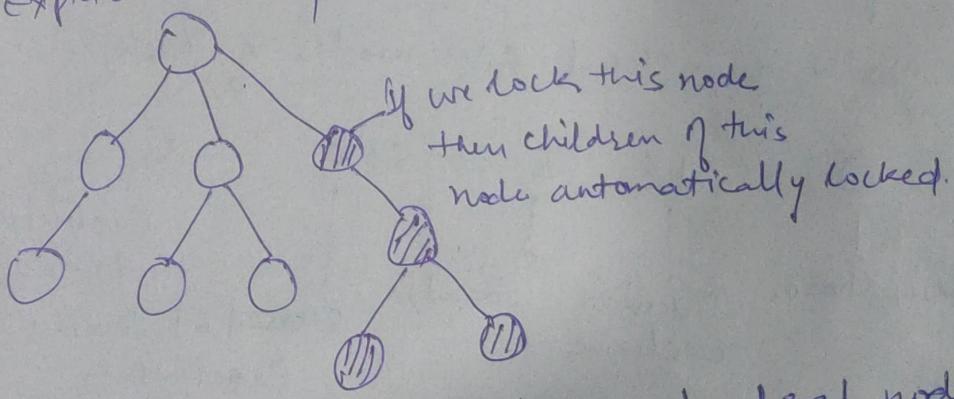
## Hierarchy of Database :-

- (1) Database
- (2) Tables
- (3) Attribute
- (4) Tuples

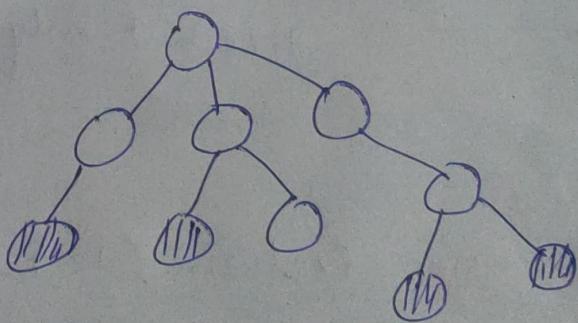
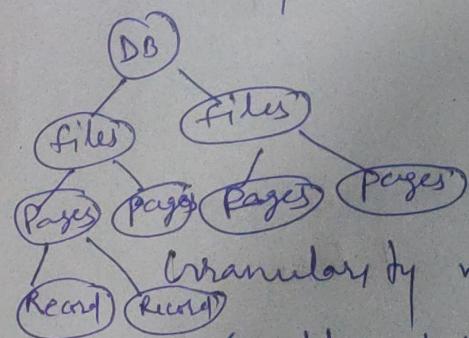
How to ~~concurrently~~ maintain concurrency control with the help of granularity.

Lock Concept :- Two types of lock:- Explicit & Implicit

### Explicit locking



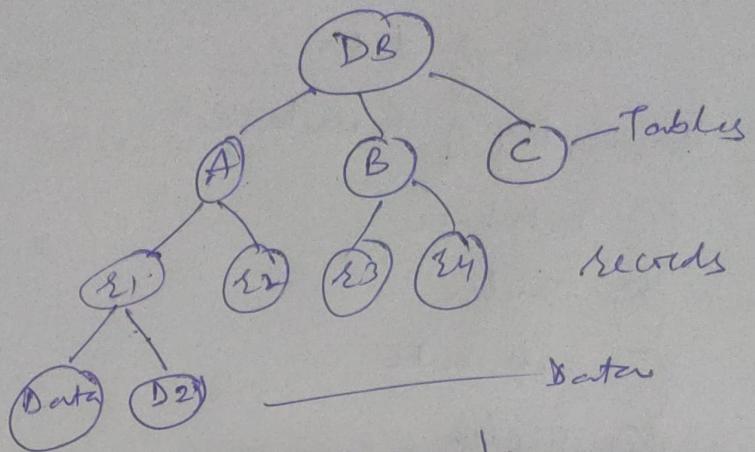
Implicit Locking :- we lock only leaf nodes in implicit locking.



Granularity means different level of data where smaller data items are listed inside bigger.

We use this granularity to ensure concurrency. We have various levels of data sizes.

→ Data items of different size. We represent this in the form of tree structure.



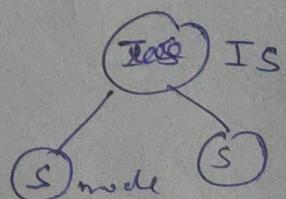
Course Granularity :- refers to large data item e.g. entire relation or a database.

Fine Granularity :- Refers to a small data items e.g. - tuples or attributes.

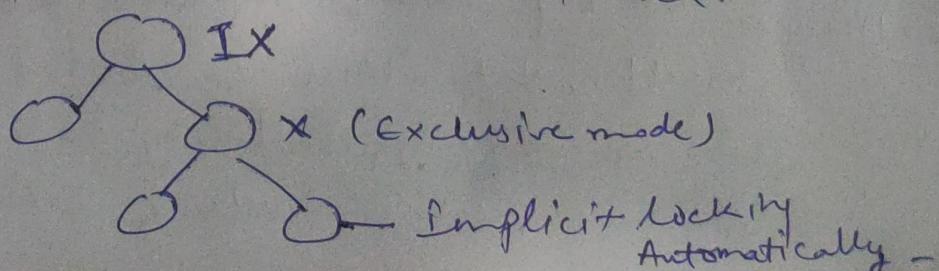
Lock Modes :- Total five modes of lock.

	IS	IX	S	SIX	X
Intension Share.					Exclusive (Read twice)
	Intension Exclusive		Share (Read)		Shared & Intension Exclusive

① Intension Share (IS) :- explicit locking at lower level of tree but only with shared lock.

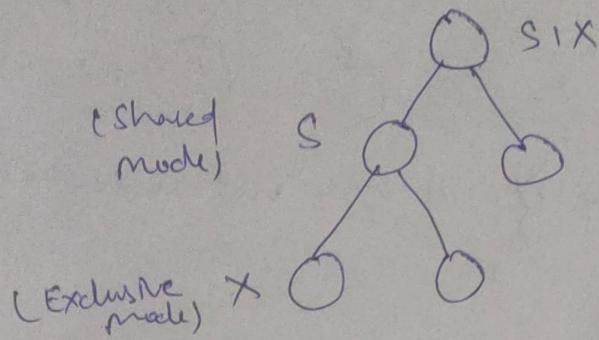


②. Intension Exclusive (IX) - explicit locking at lower level with exclusive or shared lock.



### ③ Shared & Intension Exclusive (SIX) -

Subtree rooted by that node is locked explicitly in shared mode & explicit locking is done at lower level with exclusive mode.

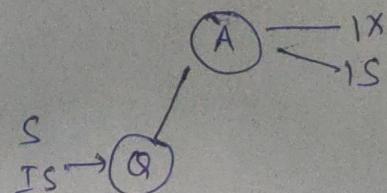


Compatibility Matrix :-

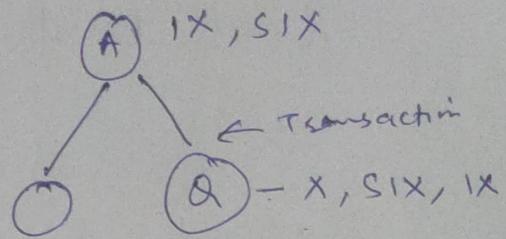
	IS	IX	S	SIX	X
IS	✓	✓	✓	✓	✗
IX	✓	✓	✗	✗	✗
S	✓	✗	✓	✗	✗
SIX	✓	✗	✗	✗	✗
X	✗	✗	✗	✗	✗

Rules:- So a transaction  $T_i$  that attempts to lock a node  $Q$  must follow these rules:-

- ① must observe the lock compatibility
- ② Root of the tree must be locked first, and it can lock in any mode.
- ③ A node  $Q$  can be locked by  $T_i$  in S OR IS mode only if the parent of  $Q$  is current locked by  $T_i$  in either IX OR IS mode.



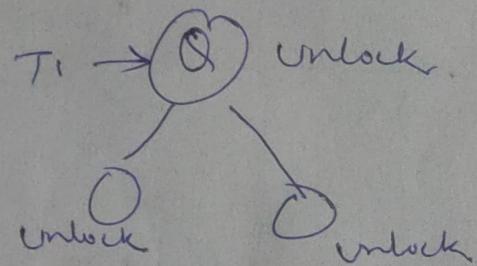
④ A node  $Q$  can be locked by  $T_i$  in X, SIX or IX mode only if the parent of  $Q$  is currently locked by  $T_i$  in either IX or SIX mode.



- ⑤  $T_i$  can lock a node only if it has not unlocked any node so far.
- ⑥  $T_i$  can unlock a node  $Q$  only if none of the child of  $Q$  is currently locked by  $T_i$ .

Locking: Root to leaf

unlocking: Leaf to Root

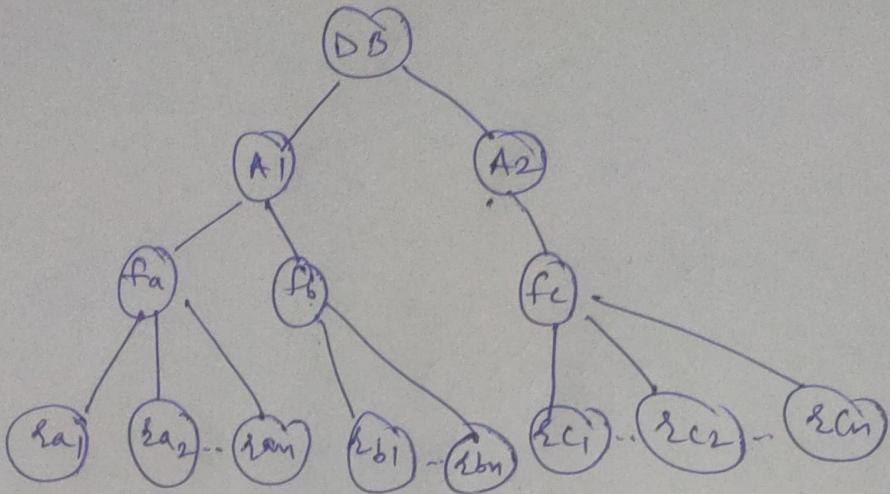


Thomas write Rule:-

$T_1$	$T_2$
	$w(Q)$
Reject * $w(Q)$ only reject	

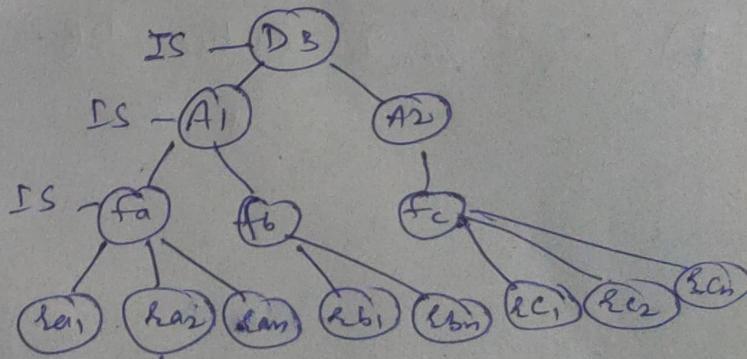
$Q = \text{to be } 30$

Ex-1

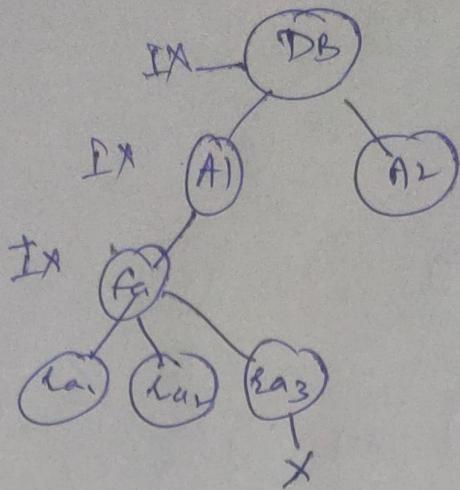


- a) T<sub>1</sub> reads ra<sub>2</sub> (DB, A1, fa  $\rightarrow$  IS || ra<sub>2</sub>  $\rightarrow$  S)
- b) T<sub>2</sub> modifies ra<sub>2</sub> X Not allocated exclusive
- c) T<sub>3</sub> modifies ra<sub>3</sub> —
- d) T<sub>4</sub> reads entire DB.
- ↳ check for compatible  
↳ feasible  
↳ feasible

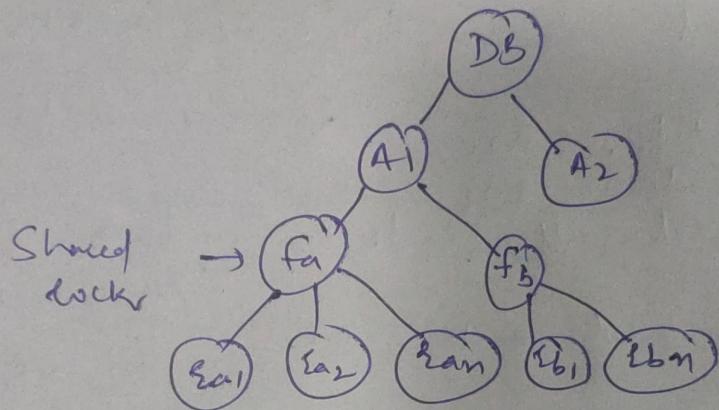
① T<sub>1</sub> reads ra<sub>2</sub> - ra<sub>2</sub> wants to read - To obtain a shared lock. But we know it travels from Root to Node. It starts from DB.



- ② T<sub>2</sub> modifies ra<sub>2</sub> - ra<sub>2</sub> already locked with Shared lock. Modifies ra<sub>2</sub> means exclusive lock. Exclusive lock will not compatible with shared lock. It won't be obtained. Transaction will rollback.
- ③ Assumes ra<sub>3</sub> - write ra<sub>3</sub> - exclusive lock, check at upper level.. compatibility at upper level. It is not compatible with X: IX lock is compatible



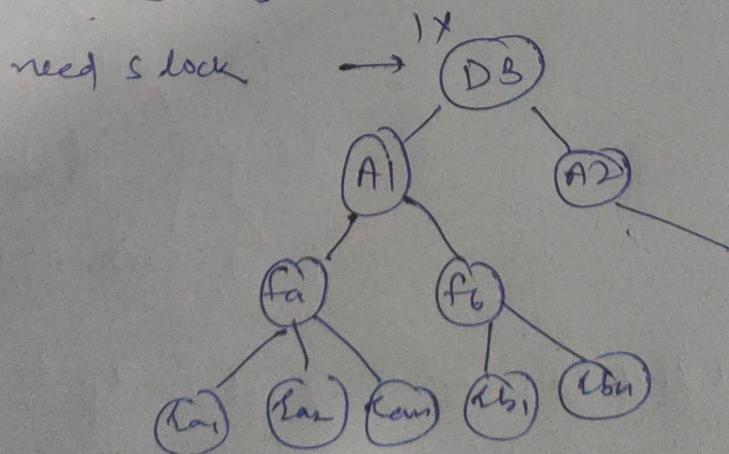
- (c) T<sub>3</sub> reads all records on fa.



Shared lock is not compatible with IX.

Shared lock can not be granted here because it is already granted IX lock.

- (d) T<sub>4</sub> reads entire DB.



Shared lock is not compatible with IX.  
We can not grant shared lock on DB.